

Monte Carlo Tree Search & Temporal-Difference

Alina Vereshchaka

CSE4/510 Reinforcement Learning
Fall 2019

avereshc@buffalo.edu

September 12, 2019

*Slides are based on Monte Carlo Tree Search, MIT 16.412J / 6.834J Cognitive Robotics
Deep Reinforcement Learning and Control, CMU 10703, Carnegie-Mellon University

Overview


- 1 Recap: Monte Carlo
- 2 Monte Carlo Tree Search
- 3 Exploration vs Exploitation
- 4 Temporal Difference

Table of Contents

- 1 Recap: Monte Carlo
- 2 Monte Carlo Tree Search
- 3 Exploration vs Exploitation
- 4 Temporal Difference

Recap: Monte-Carlo Control

$$\pi_0 \xrightarrow{E} q_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} q_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} q_*$$



- ▶ **MC policy iteration step:** Policy evaluation using MC methods followed by policy improvement
- ▶ **Policy improvement step:** greedify with respect to value (or action-value) function

Recap: Monte-Carlo Algorithm

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$\pi(s) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

Fixed point is optimal
policy π^*

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow$ return following the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$


For each s in the episode:

$\pi(s) \leftarrow \arg\max_a Q(s, a)$

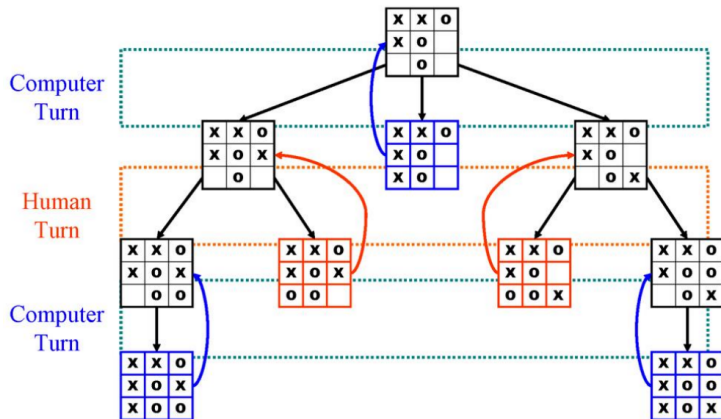
Table of Contents

- 1 Recap: Monte Carlo
- 2 Monte Carlo Tree Search
- 3 Exploration vs Exploitation
- 4 Temporal Difference

Minimize the maximum possible loss




Minimax



MCTS Outline

1. Descend through the tree
 2. Create new node
 3. Simulate
 4. Update the tree
- Repeat!

5. When you're out of time,
Return "best" child.




What do we store?

For game state k :


n_k = # games played involving k

$w_{k,p}$ = # games won (by player p)
that involved k



1. Descending

We want to **expand**,
but also to **explore**.




© Zach Weinersmith. All rights reserved. This content is excluded from our Creative Commons license. For more information, see <https://ocw.mit.edu/help/faq-fair-use/>.

Solution: *Upper Confidence Bound*

$$UCB1(k, p) = \underbrace{E[win|k, p]}_{\text{expand}} + C \underbrace{\sqrt{\frac{2 \ln(n_{parent(k)})}{n_k}}}_{\text{explore}}$$

$$\approx \frac{w_{k,p}}{n_k} + C \sqrt{\frac{2 \ln(n_{parent(k)})}{n_k}}$$

At each step,
maximize $UCB1(k, p)$



2. Expanding

Not very complicated.

Make a new node!

Set $n_k = 0$, $w_k = 0$






3. Simulating

Simulating a real game is hard.

Let's just play the game out randomly!

If we win, $\Delta = +1$. If we lose or tie, $\Delta = 0$.




4. Updating the Tree

Propagate recursively up the parents.

Given simulation result Δ ,
for each k :

$$n_{k\text{-new}} = n_{k\text{-old}} + 1$$

$$w_{k,1\text{-new}} = w_{k,1\text{-old}} + \Delta$$




5. Terminating


Return the best-ranked first ancestor!

What determines “best”?

- Highest $E[\text{win} | k]$
- Highest $E[\text{win} | k]$ AND most visited




Selection




MCTS: Expansion

Selection




Expansion




MCTS: Simulation


Selection



Expansion




Simulation




MCTS: Back-propagation


Selection




Expansion



Simulation



Backpropagation



Advantages:

- 1 Grows tree asymmetrically, balancing expansion and exploration
- 2 Depends only on the rules
- 3 Easy to adapt to new games
- 4 Heuristics not required, but can also be integrated
- 5 Complete: guaranteed to find a solution given time

Disadvantages:

- 1

Table of Contents

- 1 Recap: Monte Carlo
- 2 Monte Carlo Tree Search
- 3 Exploration vs Exploitation**
- 4 Temporal Difference


Exploration vs Exploitation

Online decision-making involves a fundamental choice:

- **Exploitation:** Make the best decision given current information (greedy)
- **Exploration:** Gather more information

The greedy algorithm selects action with highest value:

$$a_t^* = \arg \max_a Q_t(s, a)$$



Exploration vs Exploitation

ϵ – *greedy* algorithm:

- With probability ϵ choose a random action a
- With probability $1 - \epsilon$ choose “greedy” action a with the highest Q-value.

Exploration vs Exploitation

In ϵ -greedy action selection, for the case of two actions $[a_1, a_2]$ and $\epsilon = 0.5$, what is the probability that the greedy action is selected?

Table of Contents

- 1 Recap: Monte Carlo
- 2 Monte Carlo Tree Search
- 3 Exploration vs Exploitation
- 4 Temporal Difference

Monte Carlo (MC) and Temporal Difference (TD) Learning

- ▶ **Goal:** learn $v_\pi(s)$ from episodes of experience under policy π

- ▶ Incremental **every-visit Monte-Carlo**:

- Update value $V(S_t)$ toward actual return G_t

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

- ▶ Simplest **Temporal-Difference** learning algorithm: TD(0)

- Update value $V(S_t)$ toward estimated returns $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

- ▶ $R_{t+1} + \gamma V(S_{t+1})$ is called the **TD target**
- ▶ $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called the **TD error**.

DP vs. MC vs TD Learning

► Remember:

MC: sample average return
approximates expectation


$$\begin{aligned}v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] \\&= \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s\right] \\&= \mathbb{E}_{\pi}\left[R_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k R_{t+k+2} \mid S_t = s\right] \\&= \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s].\end{aligned}$$

TD: combine both: Sample
expected values and use a
current estimate $V(S_{t+1})$ of the
true $v_{\pi}(S_{t+1})$


DP: the expected values are
provided by a model. But we
use a current estimate $V(S_{t+1})$
of the true $v_{\pi}(S_{t+1})$

Dynamic Programming

$$V(S_t) \leftarrow E_\pi[R_{t+1} + \gamma V(S_{t+1})] = \sum_a \pi(a|S_t) \sum_{s', r} p(s', r|S_t, a)[r + \gamma V(s')]$$




$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$



TD(0) Method

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



TD Methods Bootstrap and Sample

- **Bootstrapping:** update involves an estimate
 - MC does not bootstrap
 - DP bootstrap
 - TD bootstrap
- **Sampling:** update does not involve an expected value
 - MC samples
 - DP does not sample
 - TD samples