

Model-Free Temporal-Difference (Q-Learning)

Alina Vereshchaka

CSE4/510 Reinforcement Learning
Fall 2019

avereshc@buffalo.edu

September 19, 2019

Overview

- 1 Recap
- 2 SARSA
- 3 Q-Learning

Table of Contents

- 1 Recap
- 2 SARSA
- 3 Q-Learning

Recap: Monte Carlo (MC) and Temporal Difference (TD) Learning

- ▶ **Goal:** learn $v_\pi(s)$ from episodes of experience under policy π

- ▶ Incremental **every-visit Monte-Carlo**:

- Update value $V(S_t)$ toward actual return G_t

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

- ▶ Simplest **Temporal-Difference** learning algorithm: TD(0)

- Update value $V(S_t)$ toward estimated returns $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$


- ▶ $R_{t+1} + \gamma V(S_{t+1})$ is called the **TD target**
- ▶ $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called the **TD error**.

TD(0) Method

- ▶ **Policy Evaluation** (the prediction problem):
 - for a given policy π , compute the state-value function v_π


- ▶ **Remember:** Simple every-visit **Monte Carlo method**:

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

 **target:** the actual return after time t

- ▶ The simplest **Temporal-Difference** method TD(0):

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$


 **target:** an estimate of the return

Example: Driving Home

<i>State</i>	<i>Elapsed Time (minutes)</i>	<i>Predicted Time to Go</i>	<i>Predicted Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

Example: Driving Home

Changes recommended by
Monte Carlo methods ($\alpha=1$)



Changes recommended
by TD methods ($\alpha=1$)




Table of Contents


1 Recap

2 SARSA

3 Q-Learning

Learning an Action-Value Function

- Estimate q_π for the **current policy** π



After every transition from a nonterminal state, S_t , do this:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

If S_{t+1} is terminal, then define $Q(S_{t+1}, A_{t+1}) = 0$

Turn this into a control method by always updating the policy to be **greedy** with respect to the current estimate

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

- SARSA is an **on-policy** algorithm which means that while learning the optimal policy it uses the current estimate of the optimal policy to generate the behaviour
- SARSA converges to an **optimal policy** as long as all state-action pairs are visited an infinite number of times and the policy converges in the limit to the greedy policy ($\epsilon = \frac{1}{t}$).

Table of Contents

- 1 Recap
- 2 SARSA
- 3 Q-Learning**

Q-learning: Off-Policy TD Control

In Q-learning the learned action-value function, Q , directly approximates the optimal action-value function, independent of the policy being followed.

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left(r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right)$$

Q-Learning Algorithm

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Loop for each step of episode:

Choose A from S using policy derived from Q (e.g., ε -greedy)

Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left[\overset{\text{Target}}{\underset{\substack{\uparrow \\ \text{Immediate Reward}}}{R} + \gamma \max_a Q(S', a)} - \overset{\text{Prediction}}{\boxed{Q(S, A)}} \right]$$

$S \leftarrow S'$

until S is terminal